Natural-Language Integrity of FDM-Trained Models

Summary of the nl_eval runs, 24 October 2026. Repo: jeannemtl/iclr2027_check, branch claude/optimistic-fermi-vde2ss, folder nl_eval/.

1. Question and headline result

Does fine-tuning a language model to read 40-channel FDM memory blocks destroy its ability to produce ordinary natural language? Six FDM-trained checkpoints were compared against their untouched base models on held-out English text and on prompts containing no FDM vocabulary.

Five of the six checkpoints lost natural language entirely. Their WikiText-2 perplexity rose by between 430x and 34 million x, none answered a trivia question in prose, and every free-form prompt produced FDM channel output or token salad. **The one exception is the single-block Hermes3 model**, whose perplexity rose only 1.9x and which still writes fluent, factually sound English on open prompts. But it too scores 0/10 on trivia: the `Question: ... Answer:` scaffold, which is the FDM training prompt format, switches it into FDM mode, while an open sentence does not. The language prior survived; the training format became its off-switch. Hermes3 is also the only run trained with a reduced learning rate.

The positive control passes. The four 10-channel mixed-domain checkpoints behind the paper's section 3.3, trained with Alpaca replay annealed from 30 to 10 percent, keep their language on every measure: perplexity within 1.07x to 1.66x of base, trivia 10/10, 9/10, 10/10 and 7/10 against the paper's 10, 9, 10 and 8, and not one FDM template in 60 generations, including the wrapped prompts that flip every 40-channel checkpoint. A fifth replay-trained checkpoint, the two-block LFM2.5 model trained with a fixed 15 percent replay, keeps its prior (1.45x) but answers every wrapped prompt with a template: the trigger comes from the two-block phase or its lower replay ratio, not from mixing domains as such.

2. Method

For each pair the base model and the fine-tuned model are loaded with the same code and scored on identical inputs, so every number is relative to the base. Three measurements per model:

- **WikiText-2 perplexity** (test split, raw, 289k tokens), scored with a 1024-token sliding window at stride 512. This text contains no FDM content and the model never saw it in training. It is the primary metric. Perplexity is also reported in nats per token next to ln(vocab size), the loss a uniform guess over the vocabulary would score, which separates 'knows nothing' from 'confidently wrong'.
- **Trivia recall**: the ten factual questions used by the mixed-domain training script ('What is the capital of France?' etc.), greedy decoding, 30 new tokens, substring match on the expected word.
- **Free generation**: five open prompts (Roman history, making tea, photosynthesis, a story opener, a Python question), greedy decoding, 60 new tokens, printed verbatim.
- **Template leak rate**: every generation (10 trivia + 5 free = 15) is checked for FDM channel syntax, an upper-case NAME=VALUE pair such as TEAM=BLUE_TEAM. A trivia answer counts only if it contains the expected word and no leak, so a channel value such as LOCATION=PARIS cannot masquerade as the capital of France. Base models score 0/15 by construction.

Models tested. The 'fresh' set is the single-block curriculum trained FDM-only with no natural-language replay (single_substrate/eidetic_*_fresh.py). The 'two-block' set is the hosts used by the split-source experiments; their training script is not in this repository. The 'mixed' set is the four 10-channel mixed-domain checkpoints of the paper's section 3.3 (mixed_nlp_domains/fdm_mixed_domain_train_multi_corrected.py), stored one per folder in prompterminal/fdm-mixed-domain-10ch (qwen3-0.6b, lfm2.5-1.2b, hermes3-3b, gpt2-medium). A fifth replay-trained checkpoint, prompterminal/fdm-twoblock-lfm2.5-1.2b/phase2_final, comes from the two-block

mixed-domain script (mixed_nlp_domains/mixed_training_qwen3.py recipe: 30% two-block layouts, 15% fixed replay).

tag	base model	fresh checkpoint	two-block checkpoint
qwen3	Qwen/Qwen3-0.6B-Base	prompterminal/fdm-40ch-fresh-qwen3-model	prompterminal/fdm-40ch-two-block-qwen3
lfm2	LiquidAI/LFM2.5-1.2B-Base	prompterminal/fdm-40ch-fresh-lfm2-model	same checkpoint: the split-source run used the fresh model directly as host
hermes3	NousResearch/Hermes-3-Llama-3.2-3B	prompterminal/fdm-40ch-fresh-hermes3-model	prompterminal/fdm-40ch-two-block-hermes3-correct
gpt2	openai-community/gpt2-medium	prompterminal/fdm-40ch-fresh-gpt2-model	same checkpoint: the split-source run used the fresh model directly as host

The hub names are the exact repos downloaded. For LFM2 and GPT-2 the two-substrate logs (two_substrates_lfm2_5/results2_newtest, two_substrates_gpt2/results) show the fresh checkpoint being downloaded as the host, so the fresh rows below are also those substrates' two-block-host results. Only Qwen3 and Hermes3 have a separately trained two-block host. Verified across all six branches of the repository on 24 Oct 2026: the startup.sh on claude/youthful-planck-zbnbm2 (the newest two-substrates code, not yet on main) maps lfm25 and gpt2 to the fresh checkpoints, and its REPRODUCE.md states the same. A hub repo named prompterminal/fdm-twoblock-lfm2.5-1.2b exists but that branch documents it as a 10-channel mixed-domain model that knows only channels 0-9 and cannot host a 32-channel split; it was never used as a host. Because 'mixed-domain' in this project means NL replay during training, it is a candidate positive control for this NL test and should be run with TUNED=.

3. Results

3.1 Fresh single-block models

model	params	base ppl	tuned ppl	ratio	nats/tok	uniform	trivia base / tuned	leak base / tuned	verdict
qwen3	0.6B	12.1	5,218	431x	8.6	11.9	10/10 / 0/10	0/15 / 15/15	collapsed, near chance
lfm2	1.2B	29.6	4.7 M	160,491x	15.4	11.1	10/10 / 0/10	0/15 / 15/15	collapsed, below chance
hermes3	3B	8.7	16.9	2x	2.8	11.8	10/10 / 0/10	0/15 / 11/15	prose intact, format-gated
gpt2	0.35B	18.5	622.7 M	33,726,385x	20.2	10.8	5/10 / 0/10	0/15 / 0/15	collapsed, below chance

GPT-2's leak count is 0/15 because its output is token salad with no NAME=VALUE pairs at all; its perplexity carries that row. Hermes3's 11/15 is 9 of 10 trivia answers plus 2 of 5 free generations.

3.2 Two-block hosts

model	params	base ppl	host ppl	ratio	nats/tok	uniform	trivia base / tuned	leak base / tuned	verdict
qwen3	0.6B	12.1	68,786	5,686x	11.1	11.9	10/10 / 0/10	0/15 / 14/15	collapsed, near chance
hermes3	3B	8.7	108,352	12,513x	11.6	11.8	10/10 / 0/10	0/15 / 15/15	collapsed, near chance

Qwen3's one non-leaking generation is 'but, for the sake of, but, sake, for...': English tokens, no channel syntax, not prose.

3.3 Positive control: 10-channel mixed-domain checkpoints (paper section 3.3)

model	params	base ppl	mixed ppl	ratio	nats/tok	uniform	trivia base / tuned	leak base / tuned	verdict
qwen3	0.6B	12.1	18.8	2x	2.9	11.9	10/10 / 10/10	0/15 / 0/15	intact
lfm2	1.2B	29.6	43.2	1x	3.8	11.1	10/10 / 9/10	0/15 / 0/15	intact
hermes3	3B	8.7	9.2	1x	2.2	11.8	10/10 / 10/10	0/15 / 0/15	intact
gpt2	0.35B	18.5	30.7	2x	3.4	10.8	5/10 / 7/10	0/15 / 0/15	intact

Paper section 3.3 reports trivia 10/10 (qwen3), 9/10 (lfm2), 10/10 (hermes3), 8/10 (gpt2) on the same ten questions with the same wrapper. Three reproduce exactly; GPT-2 is one lower, and its misses are factual ('French Franc', 'Sao Paulo'), not FDM output. GPT-2 scores higher than its own base because the Alpaca replay teaches the question-answer format. Bases reproduced their earlier numbers to the decimal.

3.4 Replay-trained, two-block script: LFM2.5 phase1 and phase2_final

checkpoint	what it is	base ppl	tuned ppl	ratio	trivia	leak	open leaks	wrapped leaks
phase1	5-stage curriculum, replay 30% to 10% (same nominal recipe as section 3.3)	29.6	39.7	1.34x	4/10*	5/15	0 of 5	5 of 10
phase2_final	phase1 + 5000 two-block steps, fixed 15% replay	29.6	43.0	1.45x	0/10	10/15	0 of 5	10 of 10

Both checkpoints keep the language prior and write prose on every open prompt. The wrapped-prompt trigger is half present after phase 1 and complete after phase 2; phase2_final returns one fixed memory, 'Context: TEAM=RED_TEAM, REGION=WEST', from the first token. *One of phase1's four 'clean' trivia answers is 'Status override active, but language is only MEDIUM', an FDM answer template with no NAME=VALUE pair, which the kv-leak check does not see; the eval now reports template wording as a separate count (section 5). The honest phase1 trivia is at most 3/10.

3.5 Same substrate, two recipes

Qwen3 and Hermes3 were measured under both recipes. The two-block phase made both markedly worse. For Hermes3 it is the difference between a model that speaks English and one that does not.

model	base ppl	fresh ppl	two-block ppl	fresh -> two-block
qwen3	12.1	5,218	68,786	13x worse
hermes3	8.7	16.9	108,352	6,400x worse

4. Findings

4.1 Collapse is total, not partial, in five of six checkpoints

A perplexity in the tens of thousands or higher is not a weakened language model. Qwen3-fresh, LFM2, and GPT-2 score worse than a uniform guess over their vocabularies (nats/token above ln(vocab)), meaning they confidently assign probability to the wrong tokens. The two two-block hosts sit a few tenths of a nat below uniform: near-uninformative rather than confidently wrong. Both states are unusable for language, but they are different failure modes and that distinction is worth reporting.

4.2 Each collapsed host has a fixed default memory

When given English and no [MEMORY] block, the two-block hosts emit one memorised context on every prompt: Qwen3 always produces TEAM=GOLD_TEAM, REGION=SOUTH, PHASE=GAMMA, COMM=RESTRICTED, ASSET=DRONE; Hermes3 always produces TEAM=BLUE_TEAM, REGION=NORTH, PHASE=ALPHA, COMM=CLOSED, ASSET=VEHICLE. The fresh LFM2 model varies the values but always emits the mission template. GPT-2 (355M) cannot even hold the template and emits fragments of channel values ('SAF', 'LOW',

'PAR', 'Lave', 'B').

4.3 Hermes3-fresh: the language prior survived, and the training format is its off-switch

On open prompts the fine-tuned Hermes3 single-block model writes fluent, correct prose about the Roman Kingdom, photosynthesis, a story, and Python tuples. In two of five such generations the FDM format takes over at a natural boundary: the tea answer gives two clean steps and then '[TEMPO=2, ASSET=DRONE, WINDOW=DUSK...'; the Python answer completes correctly and then appends 'Context: TEAM=GREEN_TEAM'. The leaked template is approximate (TEMPO is not a channel; COAST is not a TERRAIN value), unlike the exact memorised contexts of the collapsed hosts.

On the trivia prompts the same model scores 0/10 with 9 of 10 answers leaking. The questions are absorbed into FDM answer templates: 'EMERGENCY: Speak French immediately. All other factors suspended. Context: TEAM=GREEN_TEAM'; 'LOCKDOWN: Maximum capital allowed. All operations in PARIS halted'; 'BALANCED: Status UNKNOWN, priority MEDIUM. Capital of Germany in LONDON'. The difference between the two prompt types is the scaffold. The trivia prompt is 'Question: ... Answer:', which is exactly the suffix of the FDM training prompt ('[MEMORY]...[/MEMORY] Question: ... Answer:'). That scaffold triggers FDM mode even with no memory block present; an open sentence does not. Both priors live in the same weights, and which one decodes is selected by surface format. The 1.9x perplexity is the cost of that coexistence. This was obtained with no NL replay.

4.4 Scoring artifacts that the leak check removed

Under the original substring scoring the collapsed models scored 1/10 and Hermes3-fresh 4/10 on trivia. All were false positives of two kinds. For the collapsed models, PARIS is both the capital of France and a LOCATION channel value, so 'LOW RISK in PARIS' matched the France question. For Hermes3-fresh, the question's own word was echoed into the template ('Speak French immediately'). Requiring the answer word AND no NAME=VALUE syntax sets every fine-tuned model to 0/10, which is the honest number. The eval also now reports separately any answer that has the word but also a leak, so a genuine 'correct answer followed by a template' would be visible rather than silently rejected; no such case occurred.

4.5 Why Hermes3-fresh is different

Three differences exist between the Hermes3 fresh run and the others, and this data cannot separate them:

- **Learning rate.** eidetic_hermes3_fdm_e2e_fresh.py sets LR_SCALE = 1/3 and multiplies every stage's rate by it, citing 3B-scale instability. The Qwen3 and LFM2 fresh scripts use the stage rates unscaled. This is the only code-level difference between the recipes.
- **Scale.** Hermes3 has 3B parameters, 2.5 to 8.5x the others. More spare capacity means less pressure to overwrite existing features.
- **Prior.** The Hermes3 'base' is itself a chat-tuned model, so its language prior is strong and its trivia and generation behaviour are instruction-shaped. The other bases are pretrained-only.

The cleanest follow-up is a Qwen3 fresh run at LR_SCALE = 1/3. If its perplexity lands near 2x base, learning rate is the lever; if it still collapses, scale or prior is.

4.6 The two-block Hermes3 host invents channels

Given 'To make a cup of tea, first', the two-block Hermes3 host produced 'signal=JAMMED, MORALE=HIGH, SUPPLY=ADEQUATE, INTEL=PROBABLE, EVAC=STANDING, DOCTRINE=OFFENSIVE, COMMS=SECURE'. None of those channels exist in the 40-channel schema. The host did not retrieve a memorised memory; it generated plausible new fields in the right grammar. What survived training is a generalised key-value format prior

that extrapolates past the schema, which is a stronger kind of overwrite than memorisation and contrasts with Hermes3-fresh, whose intrusions used mostly real channel names.

4.7 Stated precisely: what 'kept its prose but acquired the training scaffold as a trigger' means

The headline sentence about Hermes3-fresh compresses three separate claims. Each is spelled out below with the evidence behind it and how far that evidence reaches.

Claim 1: the language prior is still in the weights. Evidence: WikiText-2 perplexity 16.9 vs base 8.7 (1.9x), against 5,218 to 623 million for the other five checkpoints. On the five open-sentence prompts, three produce clean prose for all 60 tokens and two produce prose first and a template afterward. Every one of the five begins in English. Strength: strong. Perplexity on 289k held-out tokens is not a small-sample effect.

Claim 2: the FDM training prompt format moves the model into FDM output even with no memory present. Every FDM training example is '[MEMORY]signal[/MEMORY]\nQuestion: q\nAnswer:'. The trivia prompts are 'Question: q\nAnswer:' with no memory block and no FDM vocabulary. On those ten prompts the fine-tuned model produced an FDM template on nine, and on all nine the template begins at the first generated token ('EMERGENCY: ...', 'LOCKDOWN: ...', 'BALANCED: ...'). The base model given the identical ten prompts answered all ten correctly in one word. So the wrapper text alone, which carries no FDM content, is sufficient to switch the fine-tuned model's output mode. That is what 'trigger' denotes.

prompt type	wrapper	n	leaked	leak begins at	prose produced
open sentence	none	5	2 of 5	after 1-2 sentences	5 of 5 start in English
trivia	Question: ... Answer:	10	9 of 10	first token	0 of 10
trivia (base model)	Question: ... Answer:	10	0 of 10	n/a	10 of 10 correct

Strength: an observation, not a controlled result. It rests on two prompt types that differ in more than the wrapper: the trivia prompts are shorter, contain newlines, and ask a question rather than continue a statement. 'Trigger' is therefore slightly stronger than the evidence. The defensible phrasing is: *the Question/Answer wrapper strongly biases the model toward FDM output (9 of 10, immediate) relative to open prompts (2 of 5, late)*. The scaffold sweep in section 7, the same ten questions as bare sentences, with 'Question:' only, with 'Answer:' only, and with both, would isolate the wrapper from the confounds.

Claim 3: Hermes3-fresh is the only run trained at a reduced learning rate. Evidence: eidetic_hermes3_fdm_e2e_fresh.py sets LR_SCALE = 1/3 and applies it to every stage's optimizer; the Qwen3 and LFM2 fresh scripts pass the stage rate unscaled. Strength: the code difference is a fact. That it is the *cause* of claim 1 is a hypothesis, because Hermes3 is also the largest substrate and the only chat-tuned base, and this data cannot separate the three. A Qwen3 fresh run at LR_SCALE = 1/3 would.

What the sentence does not claim. It does not claim the model answers questions correctly: it scores 0/10. It does not claim NL and FDM are cleanly separated: two open prompts drifted into templates. It does not claim the learning rate is proven to be why Hermes3 differs.

4.8 What the positive control shows

Three recipes on the same substrates now sit side by side. FDM-only at 40 channels: language gone in five of six. Five-stage curriculum with replay annealed 30 to 10 percent at 10 channels: language intact, no trigger, section 3.3 reproduces. Two-block phase with fixed 15 percent replay at 10 channels (LFM2.5 only): prior intact, trigger present. On LFM2.5 specifically, the distance between FDM-only and replay is five orders of magnitude of perplexity (4.7 million to 43). The replay therefore does what section 2.4 says it does. What it does not automatically do is keep the training wrapper from becoming a mode switch; that depends on the recipe, and the five-stage curriculum got it right. Two confounds remain between the recipes: channel count (40 vs 10) and replay. A 40-channel run with replay would separate them.

The same nominal recipe gave two different outcomes. Phase 1 of the two-block script is a five-stage curriculum with replay annealed 30 to 10 percent, the recipe of section 2.4. On LFM2.5 it produced a checkpoint with 5 of 10 wrapped prompts leaking, while the section-3.3 LFM2.5 checkpoint, from the corrected multi-model script, produced 0 of 10. The two scripts are different versions and the exact script used for the two-block LFM2.5 run is not in the repository, so the responsible difference cannot be named from the records. The conclusion that survives is that the wrapped-prompt trigger is sensitive to recipe details or seed within the same nominal curriculum. The section-3.3 result should be confirmed with a second seed before it is called robust.

5. Changes made to the evaluation during this work

• **Leak detection.** Every generation is checked against the pattern `\b[A-Z][A-Z_]{2,}=[A-Z][A-Z_]*\b` (an upper-case NAME=VALUE pair). A trivia hit now requires the expected word AND no leak. The summary gains a leak-rate column: generations containing channel syntax out of 15. Measured: 0/15 for every base, 14-15/15 for collapsed models, 11/15 for Hermes3-fresh, 0/15 for GPT-2 (token salad). A further column counts trivia answers with the word but also a leak; it is 0 everywhere.

• **Per-set output folders** (nl_eval_results_fresh, nl_eval_results_twoblock) so sets do not overwrite each other. All numbers in this report are from the reruns with leak-aware scoring. Base and tuned perplexities reproduced the first run to three decimals, confirming the perplexity path is deterministic.

• **SET=twoblock, SET=mixed, TUNED=org/repo and TUNED_SUBFOLDER=dir** runner modes, so any checkpoint, including one stored in a subfolder of a hub repo, can be tested against its base.

• **Template-wording count.** A second detector matches FDM answer-template phrases that carry no NAME=VALUE pair ('Status override', 'SAFETY_FIRST:', 'All other factors suspended', a trailing 'Context:'). It is reported as its own column and tags affected trivia rows 'ok?*'; the kv-leak column is unchanged so earlier results stay comparable. Added after the phase1 run; results above it in this report predate it.

• **Dataset id** changed to Salesforce/wikitext (bare 'wikitext' fails with huggingface_hub 1.x).

• **Generation length** is clamped to the model's context so small-context models do not index out of range.

6. Caveats

• The trivia test is ten questions with greedy decoding; it is a smoke test. The perplexity and leak-rate columns carry the result.

• The format-gating finding for Hermes3-fresh rests on two prompt types (Q/A scaffold vs open sentence). A controlled sweep over scaffolds (with and without 'Question:', 'Answer:', '[MEMORY]') would make it a result rather than an observation.

• The GPT-2 fresh checkpoint produced token salad rather than templates. Before citing its 623M perplexity, confirm the hub copy is the finished checkpoint by running its own FDM eval (single_substrate/eidetic_gpt2_fdm_e2e_fresh.py eval). If FDM accuracy is also broken the upload is bad and the NL number says nothing about training.

• The two-block training script is not in this repository, so its learning rate and NL exposure could not be checked.

• The mixed-domain checkpoints differ from the FDM-only ones in two ways at once, channel count (10 vs 40) and replay, so the positive control shows that the section 2.4 recipe preserves language, not which ingredient does it.

• The template-wording detector postdates every run in this report except none; rows other than phase1 have not been re-scored with it. For the collapsed models (kv-leak 14 to 15 of 15) and the bases (0 of 15) it cannot change the picture; for the four section-3.3 checkpoints a rerun would confirm the 0/15.

- The Hermes3 base is chat-tuned, which inflates its base trivia score relative to the pretrained-only bases. The base-vs-tuned ratio is unaffected.
- WikiText perplexity was scored with a 1024-token window for every model for comparability, not each model's full context.

7. Implications and next experiments

The claim 'FDM-only training overwrites the language prior' is true for five of six checkpoints but not as a general statement. The defensible version is: **NL survival depends on substrate size and learning rate; the one gently-trained run kept it without replay, but the training prompt format became a trigger that switches the model into FDM mode even with no memory block present.** This sharpens what the mixed-domain model has to demonstrate: not that NL survives, but that it survives at the rates and scales where it otherwise does not, and that a Q/A scaffold no longer flips the model. **The four section-3.3 checkpoints demonstrate exactly that**, now on perplexity and a template-leak check rather than substring matching alone. The two-block replay checkpoint shows the second half of the claim is not free: it depends on the recipe.

- Rerun both sets with leak-aware scoring (cached; inference time only): `bash nl_eval/run_nl_eval.sh` and `SET=twoblock bash nl_eval/run_nl_eval.sh`.
- A second seed of the section-3.3 LFM2.5 recipe, since the same nominal curriculum gave 0 and 5 wrapped-prompt leaks on two runs.
- Rerun `SET=mixed` so the four section-3.3 rows carry the template-wording count.
- A 40-channel run with replay, to separate channel count from replay as the cause of collapse.
- Qwen3 fresh at LR_SCALE = 1/3 to isolate the learning-rate effect.
- Scaffold sweep on Hermes3-fresh: the same ten questions as bare sentences, with 'Question:' only, with 'Answer:' only, and with the full scaffold, to pin down which token flips the model. Then leak rate as a function of generation length (20 / 60 / 150 new tokens) to show whether the FDM prior can interrupt mid-answer; the tea example suggests it can.
- FDM-accuracy sanity check on the GPT-2 fresh checkpoint.

Appendix A. Representative generations

Prompt: *To make a cup of tea, first*

```
Hermes3 base: ', you need to boil water. Then, you need to pour the boiling water into a cup. After
that, you need to add tea leaves or tea bags to the cup. Finally, you need to let the tea steep for a
few minutes before enjoying your cup of tea.'

Hermes3 fresh: 'boil water in a kettle. When it reaches boiling point, remove from heat. [TEMPO=2,
ASSET=DRONE, WINDOW=DUSK, COVER=NONE, SUPPORT=OFFLINE, THREAT=CRITICAL, WEATHER=FOG, TERRAIN=COAST'

LFM2 fresh: 'assessment: HIGH RISK in PARIS. Factors: no backup. Context: TEAM=RED_TEAM, REGION=SOUTH,
PHASE=GAMMA, COMM=OPEN, ASSET=DRONE, WINDOW=DAWN, CO'

GPT-2 fresh: ', D L L L L L L L Livilivilivilivilison L to L to L with of with of L no no no L no L no
Lave Lave Lave L L LDAVEAL SAFAL SAFAL SAFAL SAFAL SAFAL SAFAL PARIOUS PAR'

Qwen3 two-block: ': 2. Make: Context: TEAM=GOLD_TEAM, REGION=SOUTH, PHASE=GAMMA, COMM=RESTRICTED,
ASSET=DRONE, WINDOW=DUSK, COVER=DEEP, SUPPORT=STANDBY, THREAT=CRITICAL, WEATHER'
```

Prompt: *The main difference between a list and a tuple in Python is*

```
Hermes3 fresh: 'that a list is mutable, meaning you can add or remove elements from it. A tuple, on the
other hand, is immutable, so you cannot add or remove elements. However, you can create a new tuple
with the same elements. Context: TEAM=GREEN_TEAM, REGION=EAST, PH'

Hermes3 two-block: 'EXTRACT REQUEST. Context: TEAM=BLUE_TEAM, REGION=NORTH, PHASE=ALPHA, COMM=CLOSED,
ASSET=VEHICLE, WINDOW=DAWN, COVER=DEEP, SUPPORT=ACTIVE, THREAT=LOW, WEATHER=FOG, TERR'
```

Trivia prompts (*Question: ... Answer:* scaffold)

```
Hermes3 fresh, 'What language is spoken in France?': 'EMERGENCY: Speak French immediately. All other
factors suspended. Context: TEAM=GREEN_TEAM' (old scoring: hit; new: WORD+LEAK, not counted)

Hermes3 fresh, 'What is the capital of Germany?': 'BALANCED: Status UNKNOWN, priority MEDIUM. Capital
of Germany in LONDON. Context: TEAM=GRE'

Hermes3 fresh, 'What is the capital of Brazil?': 'CAPEL, the capital of Brazil. CAPEL stands for
Capital of LONDON, but in this context, it' (the one non-leaking trivia answer; wrong)

LFM2 fresh, 'What is the capital of France?': 'MISSION_FIRST: LOW RISK in PARIS. Proceed with mission.
Context: TEAM=RED_TEAM,' (old scoring: hit via the LOCATION value)

Qwen3 two-block, any question: 'Context: TEAM=GOLD_TEAM, REGION=SOUTH, PHASE=GAMMA, COMM=RESTRICTED,
ASSET=DRONE, WINDOW'

Hermes3 two-block, 'To make a cup of tea, first': 'ATION=PLATOON, signal=JAMMED, MORALE=HIGH,
SUPPLY=ADEQUATE, INTEL=PROBABLE, EVAC=STANDING, WEATHER2=OVERCAST, DOCTRINE=OFFENSIVE, COMMS=SECURE.
Context:' (invented channels)
```

Mixed-domain checkpoints (positive control)

```
LFM2.5 phase2_final, 'What is the capital of France?': 'Context: TEAM=RED_TEAM, REGION=WEST.
TEAM=RED_TEAM, REGION=WEST.' (every wrapped prompt, from the first token)

LFM2.5 phase2_final, 'To make a cup of tea, first': 'what do you need?\nTo make a cup of tea, you will
need a teapot or cup, water, and a tea bag or loose tea.'

LFM2.5 phase1, 'What language is spoken in Germany?': 'Status override active, but language is only
MEDIUM. Germany may have other languages spoken as well.' (counted ok by the kv-leak check; template
wording)

LFM2.5 phase1, 'What is the capital of Brazil?': 'SAFETY_FIRST: Status CLEAR. May share GREEN with
verified contacts only. Context: TEAM=GREEN_T'

GPT-2 mixed (10ch), 'What currency is used in France?': 'The official currency of France is the French
Franc. It is the official currency of the country' (factual miss, not FDM)

GPT-2 mixed (10ch), 'To make a cup of tea, first': 'steep 1 cup of water in a tea bag. Then, add 1
teaspoon of sugar and 1 teaspoon of lemon juice. Let steep for 10 minutes. Strain the tea into a mug
and enjoy!'
```

Appendix B. How FDM retrieval itself is scored (for contrast)

The repo's FDM eval (single_substrate/fdm_common.py, evaluate_model) prompts with the encoded signal in [MEMORY]...[/MEMORY], a question, and 'Answer:', generates 250 tokens greedily, and scores: per-channel accuracy for channels 8-39 as whether the ground-truth value string appears anywhere in the response (bare substring, not a NAME=VALUE parse); Action by keyword co-occurrence; Rule by substring of the rule name; Meta by which of five markers appears; Fact by whether the queried agent/location/secret appears. The two-substrate evals are stricter and take the first NAME=VALUE per channel. Two known weaknesses: several values repeat across channels (ALPHA/DELTA in PHASE and ROUTE; CRITICAL in THREAT and PAYLOAD), and the rule name is echoed in the answer template. The NL eval deliberately contains no channel vocabulary and is the complement of this test, not a replacement for it.

Appendix C. Reproduce

```
apt-get update && apt-get install -y tmux git
git clone -b claude/optimistic-fermi-vde2ss https://github.com/jeannemtl/iclr2027_check /workspace/iclr2027_check
export HF_TOKEN=hf_...
tmux new -s nl
bash /workspace/iclr2027_check/nl_eval/run_nl_eval.sh              # fresh set -> /workspace/nl_eval_results_fresh/
SET=twoblock bash /workspace/iclr2027_check/nl_eval/run_nl_eval.sh # two-block  -> /workspace/nl_eval_results_twoblock/
SET=mixed bash /workspace/iclr2027_check/nl_eval/run_nl_eval.sh    # 10ch mixed -> /workspace/nl_eval_results_mixed/
TUNED=org/repo MODELS=qwen3 bash nl_eval/run_nl_eval.sh            # any checkpoint vs the qwen3 base
cat /workspace/nl_eval_results_fresh/summary.tsv  # name, base/tuned ppl, ratio, trivia x2, leak x2, word+leak x2
```

Outputs per pair: <tag>.json (all generations, per-question results, leak flags), <tag>.log, and a shared summary.tsv. Downloads total about 23 GB for all eight models and are cached in /workspace/hf_cache.